

# Introduction to Data Management

## Joining Tables

Paul G. Allen School of Computer Science and Engineering  
University of Washington, Seattle

# Outline

- Introduce Joins
- Demo in Sqlite

# A little extra SQL

- ORDER BY – Orders result tuples by specified attributes (default ascending)

```
SELECT P.UserID, P.Name, P.Salary  
FROM Payroll AS P  
ORDER BY P.Name ASC
```



Default

```
SELECT P.UserID, P.Name, P.Salary  
FROM Payroll AS P  
ORDER BY P.Salary DESC
```

# A little extra SQL

- ORDER BY – Orders result tuples by specified attributes (default ascending)

```
SELECT P.UserID, P.Name, P.Salary  
FROM Payroll AS P  
ORDER BY P.Salary, P.Name;
```

| UserID | Name    | Salary |
|--------|---------|--------|
| 123    | Jack    | 50000  |
| 345    | Allison | 50000  |
| 567    | Magda   | 90000  |
| 789    | Dan     | 100000 |



| UserID | Name    | Salary |
|--------|---------|--------|
| 345    | Allison | 50000  |
| 123    | Jack    | 50000  |
| 567    | Magda   | 90000  |
| 789    | Dan     | 100000 |

# A little extra SQL

- DISTINCT – Deduplicates result tuples
- Data exploration:  
“What are the possible jobs in this dataset?”

```
SELECT DISTINCT Job  
FROM Payroll;
```

| Job  |
|------|
| TA   |
| Prof |

# A little extra SQL

- **DISTINCT** – Deduplicates result tuples

```
SELECT P.Job  
      FROM Payroll AS P  
      WHERE P.Salary > 70000;
```

| Job  |
|------|
| Prof |
| Prof |

# A little extra SQL

- **DISTINCT** – Deduplicates result tuples

```
SELECT P.Job  
      FROM Payroll AS P  
      WHERE P.Salary > 70000;
```

| Job  |
|------|
| Prof |
| Prof |

```
SELECT DISTINCT P.Job  
      FROM Payroll AS P  
      WHERE P.Salary > 70000;
```

| Job  |
|------|
| Prof |

# A little extra SQL

- DISTINCT – Deduplicates result tuples
- Data exploration:  
“What are the possible jobs in this dataset?”



# A little extra SQL

- DISTINCT – Deduplicates result tuples
- Data exploration:  
“What are the possible jobs in this dataset?”

```
SELECT DISTINCT Job  
FROM Payroll;
```

| Job  |
|------|
| TA   |
| Prof |

# Preview!

- Data exploration:

“How many people are in this dataset?”

# Preview!

- Data exploration:

“How many people are in this dataset?”

| COUNT(*) |
|----------|
| 4        |

# Preview!

- Data exploration:

“How many people are in this dataset?”

```
SELECT COUNT (*)  
FROM Payroll;
```

| COUNT(*) |
|----------|
| 4        |

```
SELECT COUNT (*) AS num_people  
FROM Payroll;
```

to rename column

# Joins

- Foreign keys are able to *describe* a relationship between tables
- Joins are able to realize combinations of data
- Joins do **not** require a foreign key, but often they go together

# Inner Joins

- Bread and butter of SQL queries
  - “Inner join” is often interchangeable with just “join”
- Inner Join syntax:

**Payroll**

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

**Regist**

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

Join Predicate

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R ON P.UserID = R.UserID;
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
ON P.UserID = R.UserID;
```

How do we  
algorithmically  
get our results?

| Name  | Car     |
|-------|---------|
| Jack  | Charger |
| Magda | Civic   |
| Magda | Pinto   |

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Register AS R
ON P.UserID = R.UserID;
```

How do we  
algorithmically  
get our results?

| Name  | Car     |
|-------|---------|
| Jack  | Charger |
| Magda | Civic   |
| Magda | Pinto   |

Compare every possible  
combination and filter  
the results that match



# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

```
SELECT P.Name, R.Car  
  FROM Payroll AS P JOIN Regist AS R  
    ON P.UserID = R.UserID;
```


```
for each row1 in Payroll:  
  for each row2 in Regist:  
    if (row1.UserID = row2.UserID):  
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics



| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |



| Name | Car |
|------|-----|
|------|-----|

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name | Car     |
|------|---------|
| Jack | Charger |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name | Car     |
|------|---------|
| Jack | Charger |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name | Car     |
|------|---------|
| Jack | Charger |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name | Car     |
|------|---------|
| Jack | Charger |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name | Car     |
|------|---------|
| Jack | Charger |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name | Car     |
|------|---------|
| Jack | Charger |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```



# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name | Car     |
|------|---------|
| Jack | Charger |

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name | Car     |
|------|---------|
| Jack | Charger |

```
for each row1 in Payroll:
    for each row2 in Regist:
        if (row1.UserID = row2.UserID):
            output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name  | Car     |
|-------|---------|
| Jack  | Charger |
| Magda | Civic   |


```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |



| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |



| Name  | Car     |
|-------|---------|
| Jack  | Charger |
| Magda | Civic   |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name  | Car     |
|-------|---------|
| Jack  | Charger |
| Magda | Civic   |
| Magda | Pinto   |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name  | Car     |
|-------|---------|
| Jack  | Charger |
| Magda | Civic   |
| Magda | Pinto   |


```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |



| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |



| Name  | Car     |
|-------|---------|
| Jack  | Charger |
| Magda | Civic   |
| Magda | Pinto   |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name  | Car     |
|-------|---------|
| Jack  | Charger |
| Magda | Civic   |
| Magda | Pinto   |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```



# Nested-Loop Semantics

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name  | Car     |
|-------|---------|
| Jack  | Charger |
| Magda | Civic   |
| Magda | Pinto   |

```
for each row1 in Payroll:
  for each row2 in Regist:
    if (row1.UserID = row2.UserID):
      output (row1.Name, row2.Car)
```

# Inner Joins

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

Explicit

```
SELECT P.Name, R.Car
FROM Payroll AS P JOIN Regist AS R
ON P.UserID = R.UserID;
```

Implicit

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID;
```

Both of them have the same meaning (for inner joins)

# Inner Joins

```
SELECT P.Name, R.Car  
      FROM Payroll AS P, Regist AS R  
      WHERE P.UserID = R.UserID;
```

- What if we have no join predicate?

```
SELECT P.Name, R.Car  
      FROM Payroll AS P, Regist AS R
```

```
for each row1 in Payroll:  
    for each row2 in Regist:  
        output (row1.Name, row2.Car)
```

- Output every possible pair: “Cross product”

# Outer Joins

Now I want to include everyone, even if they don't drive.

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

# Outer Joins

Now I want to include everyone, even if they don't drive.

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

```
SELECT P.Name, R.Car
FROM Payroll AS P LEFT OUTER JOIN Regist AS R
ON P.UserID = R.UserID;
```

# Outer Joins

Now I want to include everyone, even if they don't drive.

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

```
SELECT P.Name, R.Car
FROM Payroll AS P LEFT OUTER JOIN Regist AS R
ON P.UserID = R.UserID;
```

# Outer Joins

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name    | Car     |
|---------|---------|
| Jack    | Charger |
| Allison | NULL    |
| Magda   | Civic   |
| Magda   | Pinto   |
| Dan     | NULL    |

# Outer Joins

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

| Name    | Car     |
|---------|---------|
| Jack    | Charger |
| Allison | NULL    |
| Magda   | Civic   |
| Magda   | Pinto   |
| Dan     | NULL    |

NULL is a value placeholder. Depending on context, it may mean unknown, not applicable, etc.



# Outer Joins

- **LEFT OUTER JOIN**
  - All rows in left table are preserved
- **RIGHT OUTER JOIN**
  - All rows in right table are preserved
- **FULL OUTER JOIN**
  - All rows are preserved

# Self Joins

Find all people who drive a Civic and Pinto

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
        R.Car = 'Civic';
```

# Self Joins

Find all people who drive a Civic and Pinto

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
        R.Car = 'Civic' AND
        R.Car = 'Pinto';
```

Will this work?

# Self Joins

Find all people who drive a Civic and Pinto

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

```
SELECT P.Name, R.Car
  FROM Payroll AS P, Regist AS R
 WHERE P.UserID = R.UserID AND
        R.Car = 'Civic' AND
        R.Car = 'Pinto';
```

Will this work?

Nope, empty set is  
returned

# Self Joins

Find all people who drive a Civic and Pinto

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |
| 789    | Pinto   |

```
SELECT P.Name, R.Car
FROM Payroll AS P, Regist AS R
WHERE P.UserID = R.UserID AND
      R.Car = 'Civic' AND
      R.Car = 'Pinto';
```

Discuss with the people around you how you would solve this.

# Self Joins

Find all people who drive a Civic and Pinto

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

```
SELECT P.Name, R1.Car
FROM Payroll AS P, Regist AS R1, Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```

# Self Joins

Find all people who drive a Civic and Pinto

| UserID | Name    | Job  | Salary |
|--------|---------|------|--------|
| 123    | Jack    | TA   | 50000  |
| 345    | Allison | TA   | 60000  |
| 567    | Magda   | Prof | 90000  |
| 789    | Dan     | Prof | 100000 |

| UserID | Car     |
|--------|---------|
| 123    | Charger |
| 567    | Civic   |
| 567    | Pinto   |

All pairs of cars a person can drive

```
SELECT P.Name, R1.Car
FROM Payroll AS P, Regist AS R1, Regist AS R2
WHERE P.UserID = R1.UserID AND
      P.UserID = R2.UserID AND
      R1.Car = 'Civic' AND
      R2.Car = 'Pinto';
```

# Self Joins

- When a relation occurs twice in the FROM clause we call it a self-join



# Self Joins

- When a relation occurs twice in the FROM clause we call it a self-join
- If we have a self-join, we must use table aliases (why?)

# Self Joins

- When a relation occurs twice in the FROM clause we call it a self-join
- If we have a self-join, we must use table aliases (why?) – We will have duplicate attribute names and need to distinguish which table they are from

# Edges example

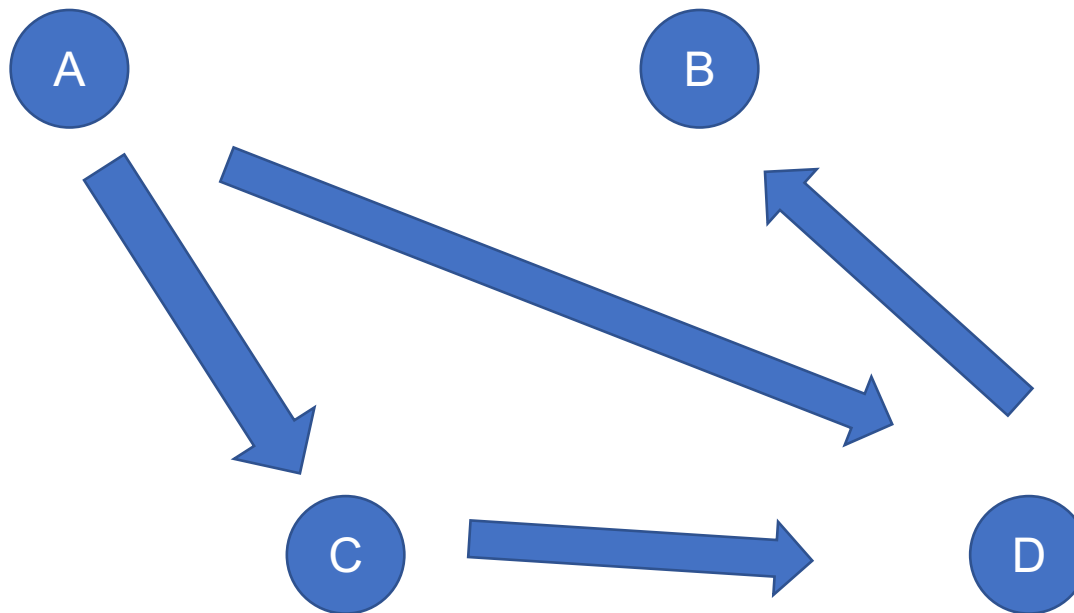
# Edges example

- Join to combine data from different tables



# Edges example

- Join to combine data from different tables

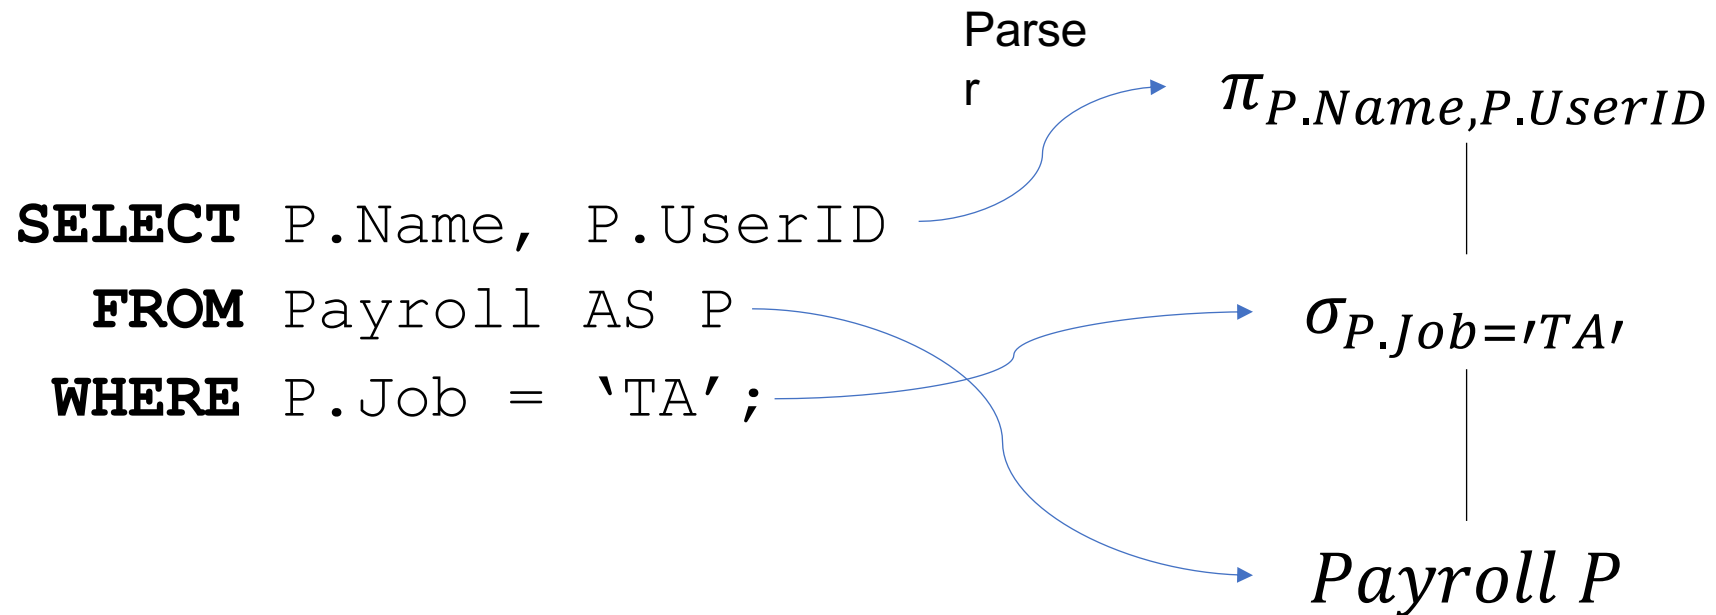


# RA Equivalencies

So far we haven't discussed equivalent RA trees.  
But all joins can be parsed directly into a “join tree”

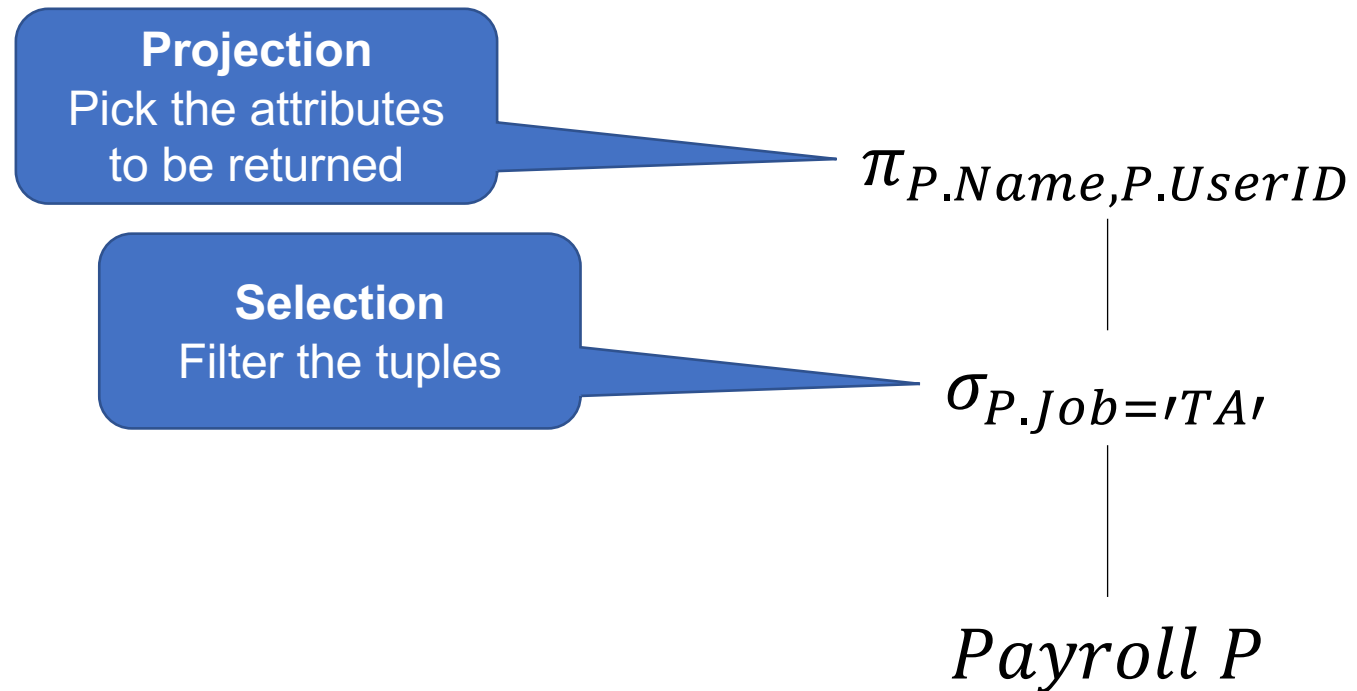
# RA Equivalencies

So far we haven't discussed equivalent RA trees.  
But all joins can be parsed directly into a “join tree”



# RA Equivalencies

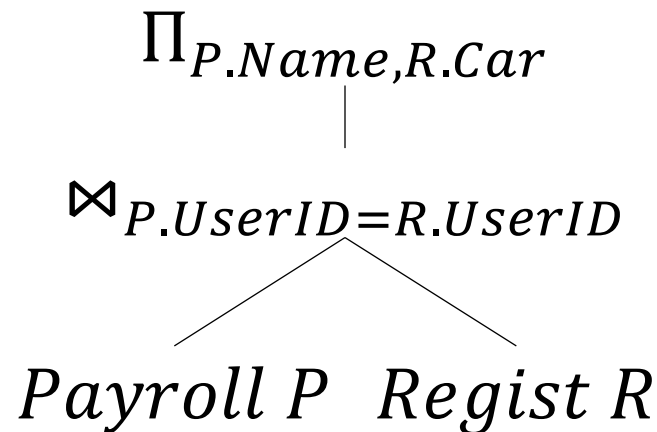
So far we haven't discussed equivalent RA trees.  
But all joins can be parsed directly into a “join tree”





# RA Equivalencies

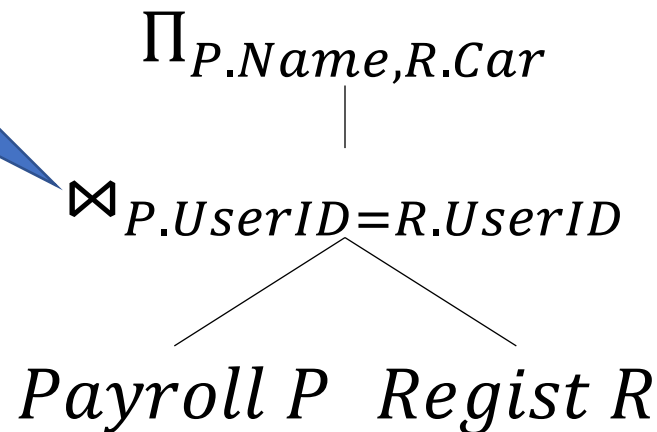
```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```



# RA Equivalencies

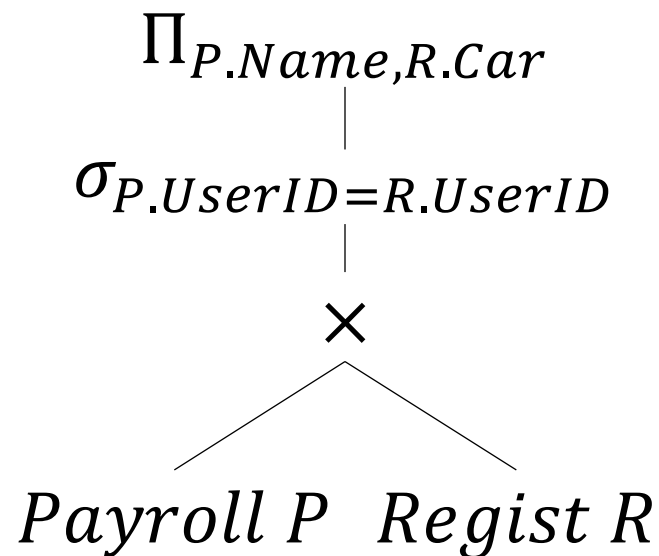
```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```

**Join**  
Combine tuples on the  
provided predicate



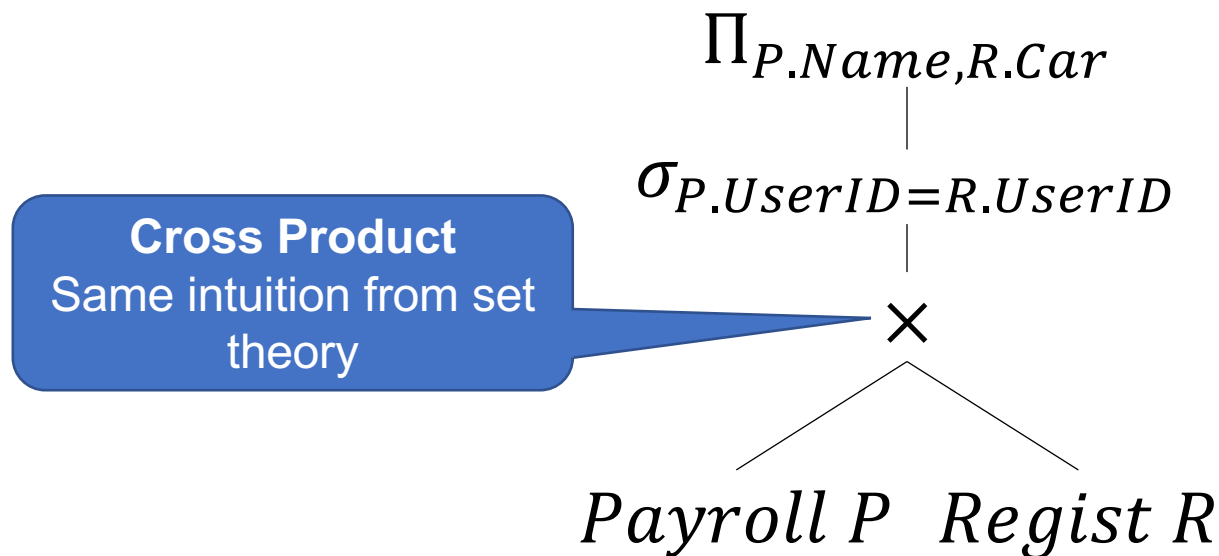
# RA Equivalencies

```
SELECT P.Name, R.Car  
  FROM Payroll AS P, Regist AS R  
 WHERE P.UserID = R.UserID;
```



# RA Equivalencies

```
SELECT P.Name, R.Car  
FROM Payroll AS P, Regist AS R  
WHERE P.UserID = R.UserID;
```



# Takeaways

- We can describe relationships between tables with keys and foreign keys
- Different joining techniques can be used to achieve particular goals
- Our SQL toolbox is growing!
  - Not just reading and filtering data anymore
  - Starting to answer complex questions